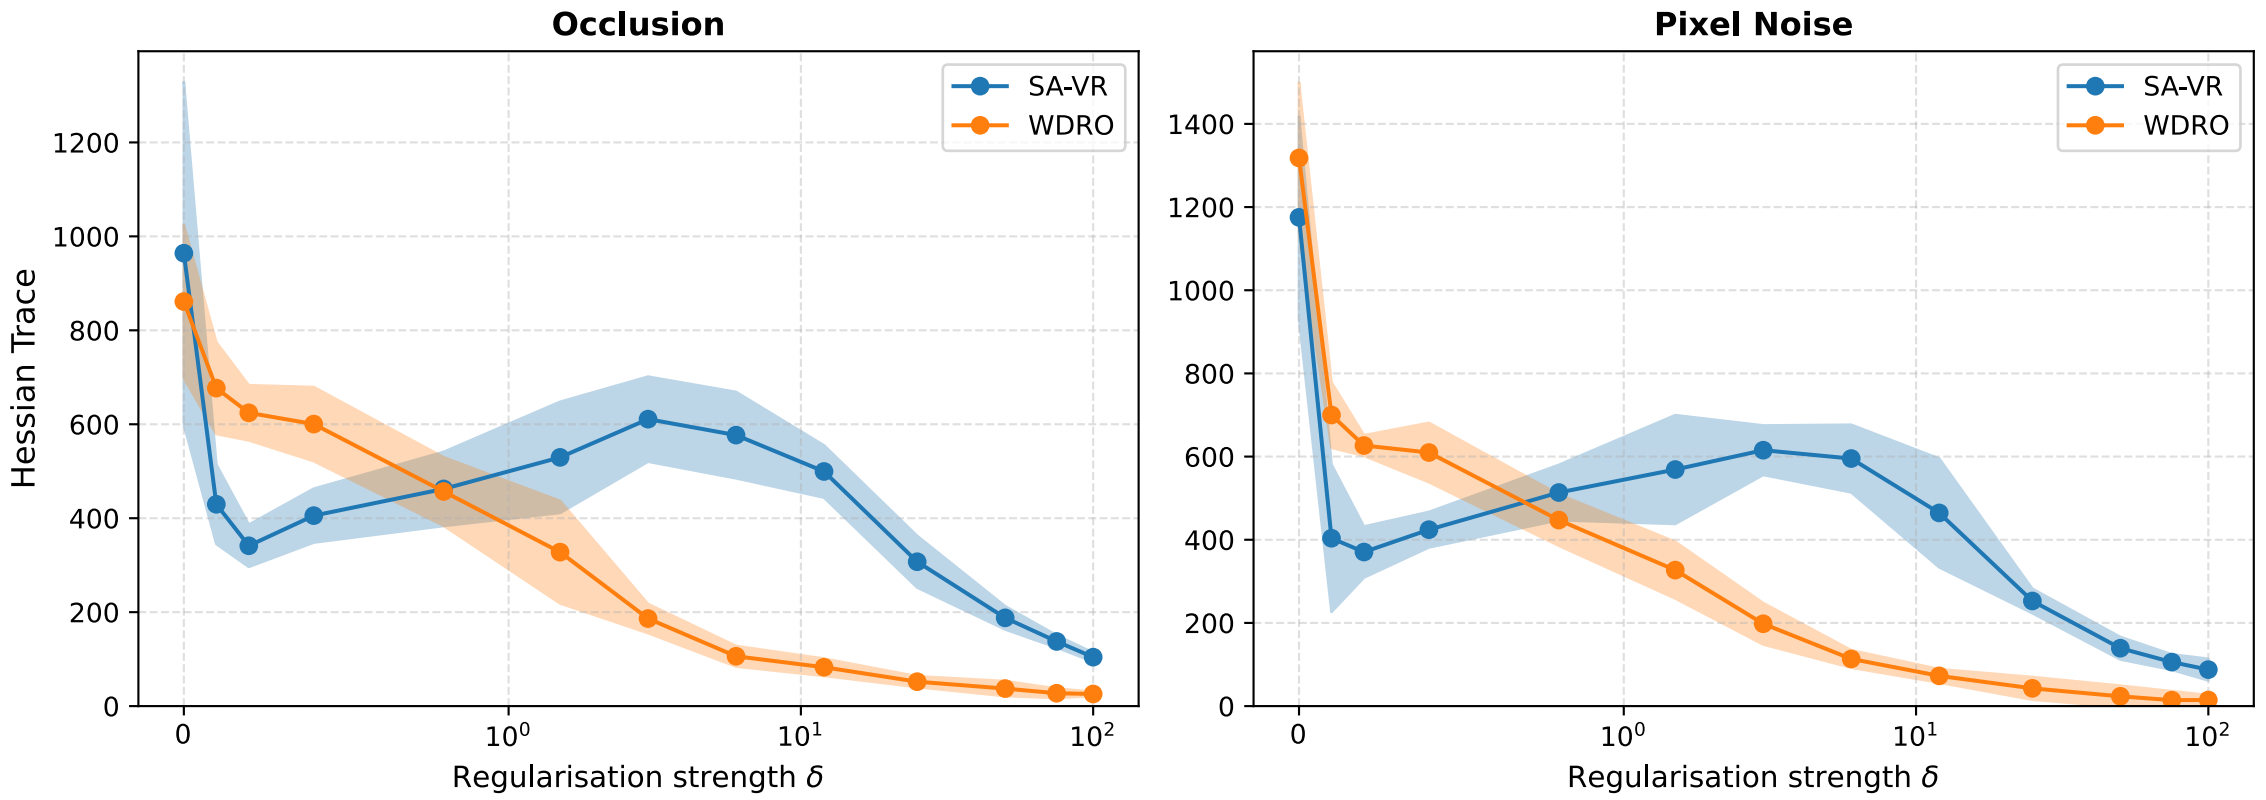
Hessian Trace — Perturbation Prob = 0.2
(lower = flatter loss landscape; CE-loss Hessian, mean ± std, 5 runs)
Occlusion
Pixel Noise
SA-VR
WDRO
Hessian Trace
Regularisation strength δ
0
200
400
600
800
1000
1200
1400
10^0
10^1
10^2

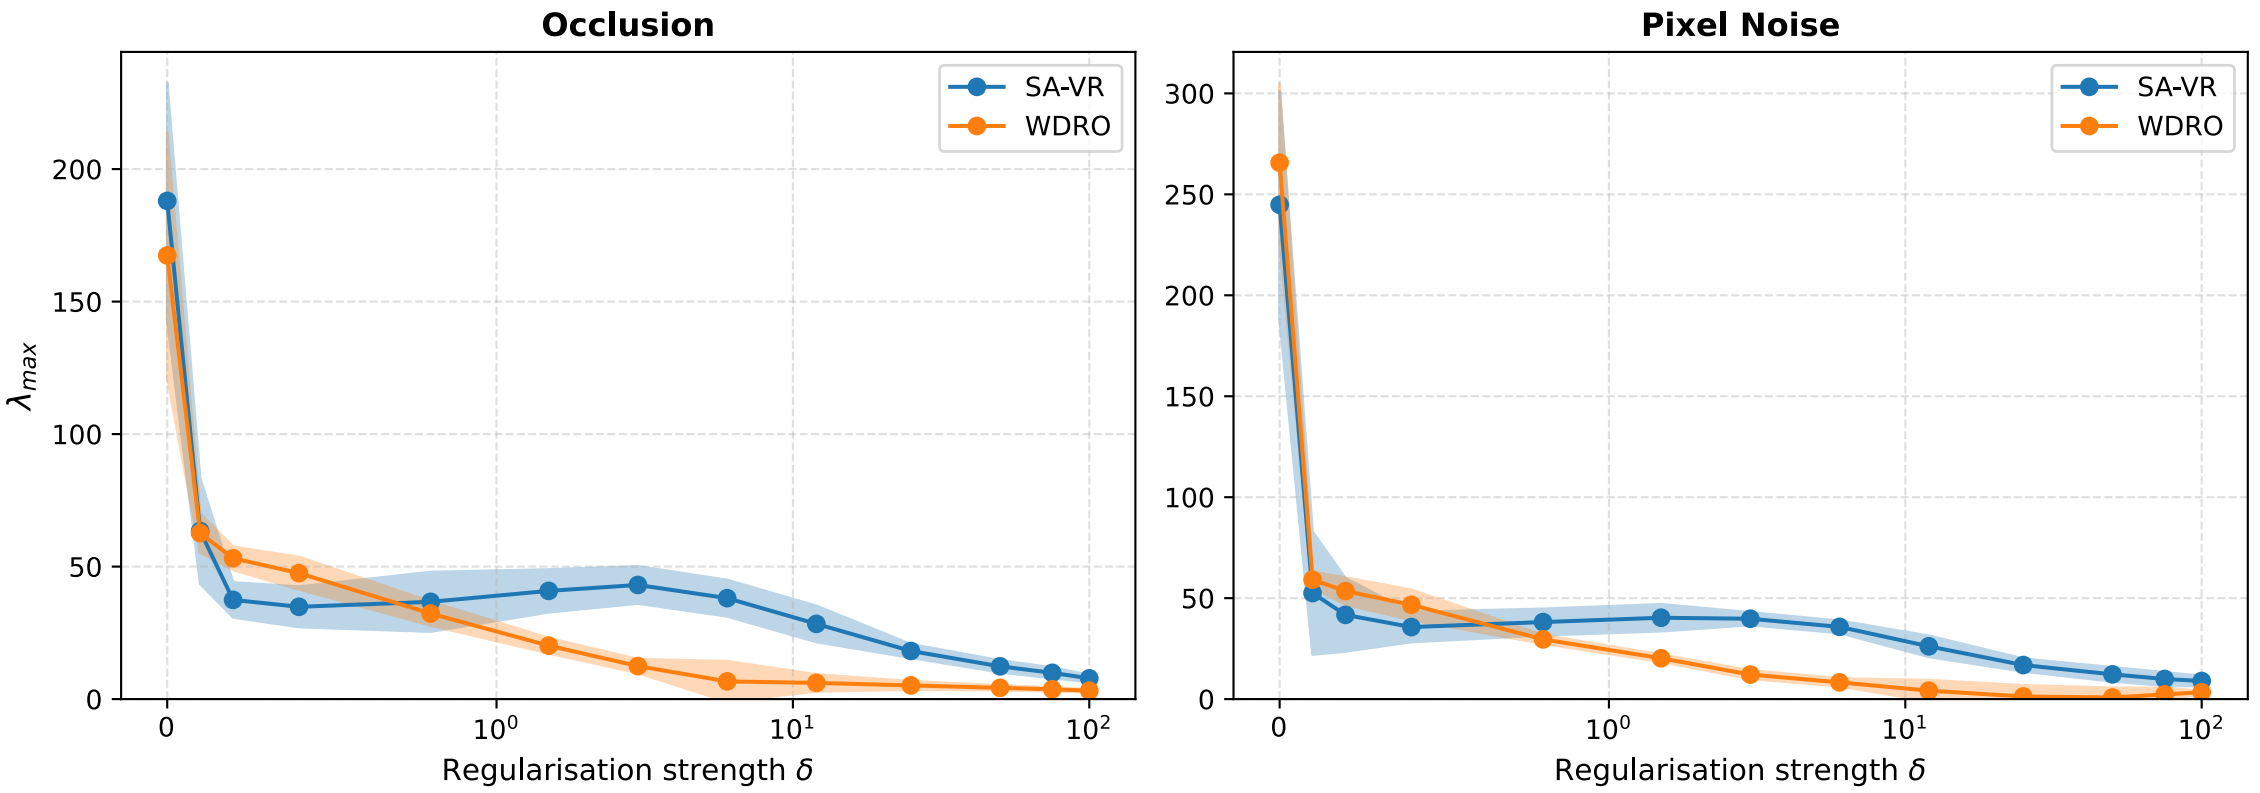
Parameter Sharpness (λ_{max}) — Perturbation Prob = 0.2
(lower = flatter minimum; CE-loss Hessian, mean ± std, 5 runs)
Occlusion
Pixel Noise
SA-VR
WDRO
λ_{max}
Regularisation strength δ
0
50
100
150
200
250
300
10^0
10^1
10^2

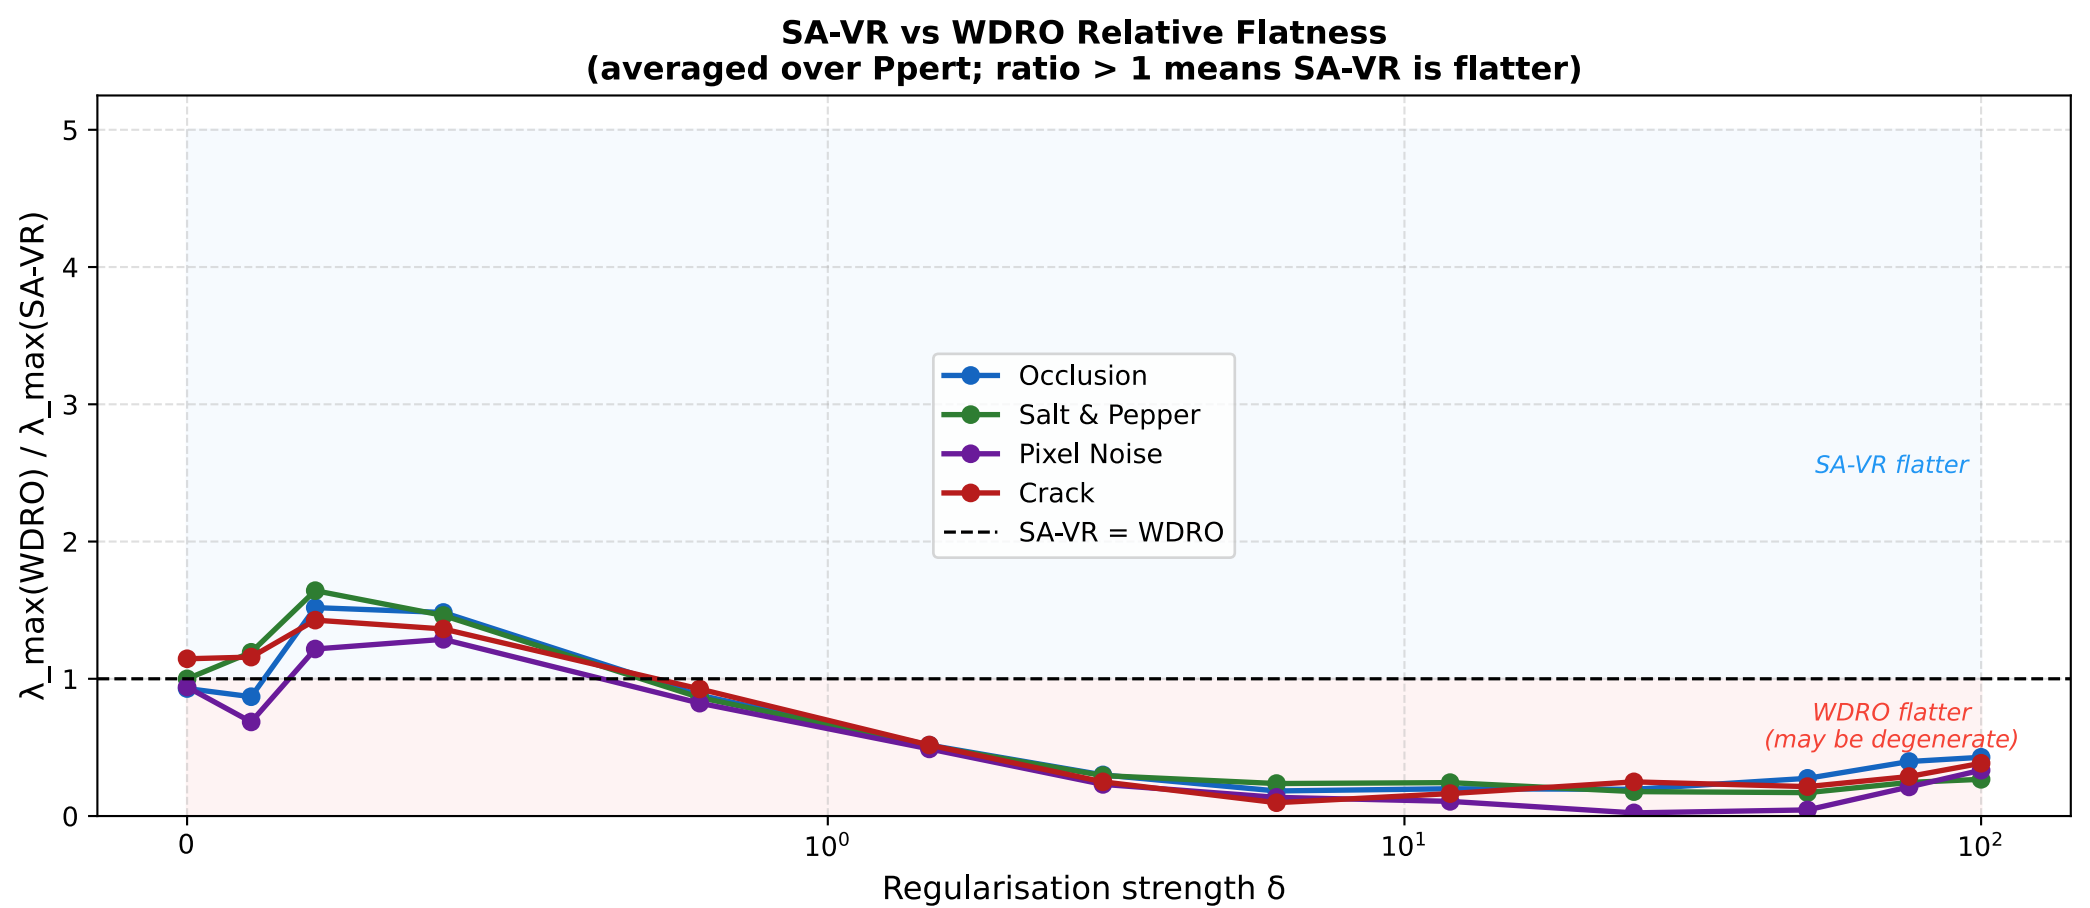
SA-VR vs WDRO Relative Flatness
(averaged over Ppert; ratio > 1 means SA-VR is flatter)
λ_max(WDRO) / λ_max(SA-VR)
Regularisation strength δ
Occlusion
Salt & Pepper
Pixel Noise
Crack
SA-VR = WDRO
SA-VR flatter
WDRO flatter
(may be degenerate)
0
1
2
3
4
5
0
10^0
10^1
10^2